

Project 4: Problem 2

20204898 박소은

Source code

```
#include <stdio.h>
#include <thrust/device_vector.h>
#include <thrust/sequence.h>
#include <thrust/transform.h>
#include <thrust/reduce.h>

#define NUM_STEPS 200000
#define STEP 1.0/NUM_STEPS

struct calculation {
    __host__ __device__
    double operator()(double i) {
        double x = (i+0.5)*STEP;
        return 4.0/(1.0+x*x);
    }
};

int main ()
{
    clock_t start_time = clock();

    // value "i" initialization
    thrust::device_vector<double> i(NUM_STEPS); // same with "i" in omp code
    thrust::sequence(i.begin(), i.end()); // 0 to NUM_STEPS in "i" vector


    thrust::device_vector<double> sum(NUM_STEPS); // vector to store the "sum"
    // same with the "for" statement in the original code
    thrust::transform(i.begin(), i.end(), sum.begin(), calculation());

    double result = thrust::reduce(sum.begin(), sum.end()); // summation
    double pi = STEP * result;

    clock_t end_time = clock();
    clock_t exe_time = end_time - start_time;
    double exe_time_sec = (double)(exe_time) / CLOCKS_PER_SEC;
    printf("Execution Time : %.10lf \n", exe_time_sec);
    printf("pi = %.10lf \n", pi);
}
```

Execution timetable & graphs

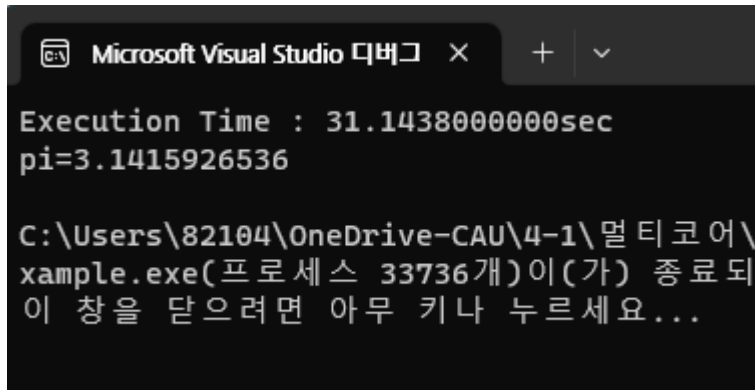
	Original Code	CUDA code
Execution 1	31.1438	0.103011
Execution 2	30.3359999	0.115084
Execution 3	34.9246633	0.094868
Execution 4	32.2134621	0.094868
Execution 5	31.2766153	0.095936
Average	31.97890812	0.1007534



Explanation/Interpretation on the results

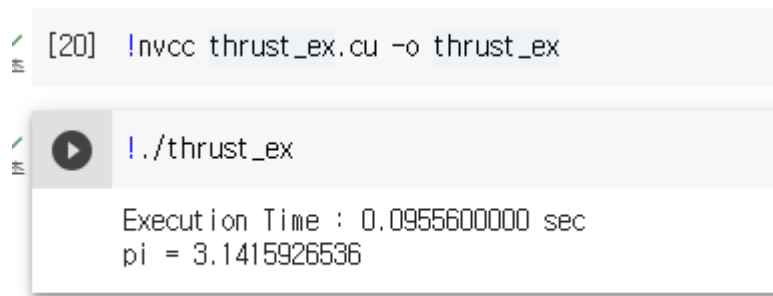
Screen captures of output results.

- Original code(omp_pi_one.c)



```
Microsoft Visual Studio 디버그 × + ▾  
Execution Time : 31.1438000000sec  
pi=3.1415926536  
  
C:\Users\82104\OneDrive-CAU\4-1\멀티코어\  
xample.exe(프로세스 33736개)이(가) 종료되  
이 창을 닫으려면 아무 키나 누르세요...
```

- thrust_ex.cu



```
[20] !nvcc thrust_ex.cu -o thrust_ex  
  
!./thrust_ex  
  
Execution Time : 0.0955600000 sec  
pi = 3.1415926536
```

Explanation / Interpretation

The average execution time of `comp_pi_one.c` is 31.97 sec, and the average execution time of CUDA code using thrust is 0.10 sec. `Thrust_ex.cu` performed about 317.39 times better than the `comp` code. This is predictable because `thrust_ex.cu` used GPU and thrust library, and the `comp_pi_one.c` was done with one thread and CPU.

Although the thread count is not directly specified in the `thrust_ex.cu` code, thrust library internally creates a CUDA kernel to perform operations. The number of threads depends on the thrust implementation and GPU, and it is processed by internally creating an optimal thread configuration.

When running this example, I used google colab's GPU T4. T4 is Tesla T4 GPU developed by NVIDIA. With T4 and Thrust, a maximum of 1024 threads can be used, but it is difficult to accurately predict the number of threads actually used. However, as can be seen from the results, it is clear that it is calculated more efficient and faster than `comp_pi_one.c`, which specifies the number of threads as 1.